

# Specialization of Run-time Configuration Space at Compile-time: An Exploratory Study

Xhevahire Tërnav<sup>1</sup>, Mathieu Acher<sup>12</sup>, Benoit Combemale<sup>1</sup>


<sup>1</sup>Université de Rennes 1, CNRS, Inria, IRISA

<sup>2</sup>Institut Universitaire de France (IUF)  
Rennes, France

SAC SE 2023, March 27 - 31, 2023

# Context


- Today's software systems are highly-configurable, e.g.,
  - ⊛ Linux:  $\approx 20K$  options;  $\approx 2^{20,000}$  configuration space (?)
  - ⊛ Renault:  $\approx 10^{22}$  cars (?) - 5 million light-years <sup>[1]</sup>
  - ⊛ x264: 39 compile-time options and 162 run-time options



1: Renault Group, Keynote Yves Bossu at the 17th International Working Conference VaMoS. 2023

# Context

- Up to 54.1% of options are rarely set by any user <sup>[1]</sup>  
→ i.e., they bloat the software
- Software debloating is the process of eliminating superfluous functionalities from it, with the goal of reducing its execution time, resource consumption, and attack surface




---

1: Xu T. et al. Hey, you have given me too many knobs!: Understanding and dealing with over-designed configuration in system software. In the 10th Joint Meeting on Foundations of SE. 2015

# Context

- Up to 54.1% of options are rarely set by any user <sup>[1]</sup>  
→ i.e., they bloat the software
- Software debloating is the process of eliminating superfluous functionalities from it, with the goal of reducing its execution time, resource consumption, and attack surface



- Debloating source code, compiled binaries, libs, or flags
- Idea: specializing the run-time configuration space of a system

1: Xu T. et al. Hey, you have given me too many knobs!: Understanding and dealing with over-designed configuration in system software. In the 10th Joint Meeting on Foundations of SE. 2015

2: Source: <https://www.educative.io/edpresso/what-is-software-debloating>

# x264: motivation case study

Most of the C-based systems are configurable through the compile-time and run-time options. For example, the system of **x264 video encoder**

**x264:** db0d417 commit; 114,475 LoC; 2:

**x264's options:** 39 compile-time; 162 run-time

E.g., x264 build without mp4 support

```
$ ./configure --disable-lsmash && make
```

E.g., video encoding with x264 using three run-time options

```
$ x264 --no-cabac --mbtree --mixed-refs -o vid.264 vid.y4m
```



**x264** has 10 presets, but they contain only 22 options, *i.e.*, 140 unchanged

# Problem statement

## Proposal: debloating run-time options at compile time

In a given context, some run-time options may never be set  
e.g., `--no-cabac` in x264

- Likely, they increase the system's binary size,
- they augment the system's attack surface (# gadgets),
- they may slow down the system, ...

Removing unused run-time variability in a software system is not trivial


Vision:

- Developers should have the means to remove unused run-time options
- And, to possibly build new x264's variants

*e.g., x264-hq, x264-fast, x264-tinyfast, x264-secured, etc.*


# Our specialization approach

## Debloating process



# Our specialization approach

## Debloating process




**T1:** Implementation patterns (getopt.h, a variable per option, etc.)



# Our specialization approach

## Debloating process




**T1:** Implementation patterns (getopt.h, a variable per option, etc.)

**T2:** Options are made explicit using C preprocessor directives

# Our specialization approach

## Debloating process




**T1:** Implementation patterns (getopt.h, a variable per option, etc.)

**T2:** Options are made explicit using C preprocessor directives

```
1  /* File removeoption.h */
2  #ifndef CABAC_YES
3  #define CABAC_YES 0
4  #endif
5  #ifndef CABAC_NO
6  #define CABAC_NO 1
7  #endif
8  /* The rest of the directives are omitted */
```

# Our specialization approach

## Debloating process



**T1:** Implementation patterns (getopt.h, a variable per option, etc.)


**T2:** Options are made explicit using C preprocessor directives

```
1  /* File encoder/encoder.c */
2  /*The previous code is omitted*/
3  if( h->param.b_cabac )
4      x264_cabac_init( h ); //option --cabac
5  else
6      x264_cavlc_init( h ); //option --no-cabac
7  /*The rest of the code is omitted*/
```

```
1  /* File encoder/encoder.c */
2  /*The previous code is omitted*/
3  #if CABAC_YES
4      if( h->param.b_cabac )
5          x264_cabac_init( h ); //option --cabac
6  #endif
7  #if CABAC_YES && CABAC_NO
8      else
9  #endif
10 #if CABAC_NO
11     x264_cavlc_init( h ); //option --no-cabac
12 #endif
13 /*The rest of the code is omitted*/
```

# Our specialization approach

## Debloating process




**T1:** Implementation patterns (getopt.h, a variable per option, etc.)

**T2:** Options are made explicit using C preprocessor directives

**T3:** Preprocessor switches

# Our specialization approach

## Debloating process



- T1:** Implementation patterns (getopt.h, a variable per option, etc.)
- T2:** Options are made explicit using C preprocessor directives
- T3:** Preprocessor switches
- T4:** The specialized system should be compilable and valid (video size, quality, and system interoperability)
  - If not **T4**, we then repeat **T1** to **T4**

# Experimental design

## Null and alternative hypothesis

$H_{01}$  : The baseline software system and its specialization are not significantly different with respect to their binary size

$H_{A1}$  : The specialized software system has smaller binary size than its baseline

$H_{02}$  : The baseline software system and its specialization are not significantly different with respect to their attack surface

$H_{A2}$  : The specialized software system has smaller attack surface than its baseline

$H_{03}$  : The baseline software system and its specialization do not perform significantly different

$H_{A3}$  : The specialized software system performs better than its baseline as for the encoding time, bitrate, and frame rate

# Experimental design

## Sample of 20 specializations

- Baseline system ( $S_0$ ): x264 with default configuration in [db0d417](#)
- 20 specialized systems: 10 run-time options and 10 presets

Options \ Presets		ultrafast	superfast	veryfast	faster	fast	medium	slow	slower	veryslow	placebo
		$S_{11}$	$S_{12}$	$S_{13}$	$S_{14}$	$S_{15}$	$S_{16}$	$S_{17}$	$S_{18}$	$S_{19}$	$S_{20}$
--no-mixed-refs	$S_1$	●	●	●	●	○	○	○	○	○	○
--no-mbtree	$S_2$	●	●	○	○	○	○	○	○	○	○
--no-cabac	$S_3$	●	○	○	○	○	○	○	○	○	○
--no-weightb	$S_4$	●	○	○	○	○	○	○	○	○	○
--no-psy	$S_5$	○	○	○	○	○	○	○	○	○	○
--mixed-refs	$S_6$	○	○	○	○	●	●	●	●	●	●
--mbtree	$S_7$	○	○	●	●	●	●	●	●	●	●
--cabac	$S_8$	○	●	●	●	●	●	●	●	●	●
--weightb	$S_9$	○	●	●	●	●	●	●	●	●	●
--psy	$S_{10}$	●	●	●	●	●	●	●	●	●	●

○ - unused option; ● - used option

Scenario<sub>1</sub> :  $S_1 - S_{10}$  system's specializations by a single option

Scenario<sub>2</sub> :  $S_{11} - S_{20}$  system's specializations regarding a preset

# Experimental design

The inputs and workstation used

**Inputs:** 8 video, representative of 1300 videos in [YouTube UCG dataset](#)

Name [.mkv]	Size [MB]	Length [sec.]	Name [.mkv]	Size [MB]	Length [sec.]
V_1_720x480	108.4	13	V_5_640x360	172.8	20
V_2_480x360	155.6	20	V_6_640x360	41.5	20
V_3_640x360	165.5	19	V_7_640x360	207.4	20
V_4_640x360	172.5	19	V_8_624x464	217.2	20

**Settings:** Workstation: Fedora 33

CPU: Intel Core i7-10610U with 15.3 GiB of memory


Compiler: GCC 10.3.1

**Repeated:** 5 times, sequentially

**Hypothesis testing:** Wilcoxon signed-rank test ( $\alpha = 0.05$ , one-sided)



## $H_{01}, H_{A1}$ : The binary size of specialized system




$S_0$  : 3,096,176 bytes;  $S_{1-20}$  :  $(3,020,417 \pm 63,641)$  bytes

$H_{01}$  is rejected in favor of  $H_{A1}$  as  $p = 9.54 \cdot 10^{-7} < \alpha = 0.05$ .

*i.e.*, specializing a software system regarding its run-time options will statistically significantly reduce its binary size.

# $H_{02}, H_{A2}$ : The attack surface of specialized system

**Gadgets:** small code sequences in a system that end in a RET or JMP return instruction




$S_0$  : 106,495 gadgets;  $S_{1-20}$  :  $(103,414 \pm 2,696)$  gadgets

$H_{02}$  is rejected in favor of  $H_{A2}$  as  $p = 3.54 \cdot 10^{-4} < \alpha = 0.05$ .  $H_{02}$  is not rejected for  $S_1 - S_{10}$  ( $p = 0.25$ ) and is rejected for  $S_{11} - S_{20}$  ( $p = 9.77 \cdot 10^{-4}$ ).

*i.e.*, it is better to specialize a software system by multiple run-time options, as it will statistically significantly reduce its attack surface.

## $H_{03}$ , $H_{A3}$ : The performance of specialized system



$H_{03}$  is rejected in favor of  $H_{A3}$  for encoding time ( $p = 0.01$ ) and frame rate ( $p = 0.02$ ), but not for bitrate ( $p = 0.16$ ).  $H_{03}$  is not rejected for  $S_1 - S_{10}$  ( $p = 0.5$ ,  $p = 0.16$ , and  $p = 0.46$ , respectively). But, it is rejected for  $S_{11} - S_{20}$  for encoding time ( $p = 1.95 \cdot 10^{-3}$ ) and frame rate ( $p = 0.02$ ).

*i.e.*, specializing a software system regarding its run-time options could significantly improve its performance (in x264, its encoding time and frame rate).

# The trade-off among the system properties

System	Binary size	Gadgets	Encoding time	Bitrate	Frame rate
$S_1$	-0.270%	-0.51%	0.89%	0.00%	-1.78%
$S_2$	-0.001%	0.18%	-5.38%	0.96%	8.53%
$S_3$	-3.254%	-4.00%	0.87%	0.00%	-0.65%
$S_4$	-0.001%	-0.02%	1.06%	0.00%	-1.88%
$S_5$	-0.001%	0.42%	1.29%	0.00%	-2.26%
$S_6$	-0.403%	0.29%	1.07%	0.00%	-2.91%
$S_7$	-0.543%	-0.81%	-4.73%	0.00%	1.86%
$S_8$	-5.416%	-6.60%	-0.14%	0.00%	2.32%
$S_9$	-0.003%	0.05%	1.58%	0.00%	2.77%
$S_{10}$	-0,001%	0.18%	NaN	NaN	NaN
$S_{11}$	-6,365%	-7.18%	1.71%	0.00%	-0.68%
$S_{12}$	-4.202%	-4.88%	-3.07%	0.00%	-0.91%
$S_{13}$	-3.659%	-4.43%	-9.58%	2.50%	8.15%
$S_{14}$	-3.659%	-4.43%	-9.58%	2.50%	8.15%
$S_{15}$	-3.526%	-4.35%	-11.97%	-0.67%	14.77%
$S_{16}$	-3.526%	-4.35%	-11.28%	-0.45%	17.64%
$S_{17}$	-3.526%	-4.35%	-7.24%	-0.98%	18.65%
$S_{18}$	-3.526%	-4.35%	-5.41%	-1.09%	13.91%
$S_{19}$	-3.526%	-4.35%	-6.59%	3.71%	12.52%
$S_{20}$	-3.526%	-4.35%	-4.25%	2.49%	8.41%
Avr.	-2.447%	-2.89%	-3.40%	0.36%	5.47%

# Specialization of Run-time Configuration Space at Compile-time: An Exploratory Study

Xhevahire Tërnavá, Mathieu Acher, and Benoit Combemale

## Contributions:

- A semi-automated and sound system specialization approach
- Quantify benefits: binary size (2.45%), attack surface (2.89%), encoding time (3.40%), bitrate (0.36%), and frame rate (5.47%)
- 4/5 analysed properties show a significant improvement
- Design of the study and experimental protocol
- An available ground truth and data for reproduction and replication:  
<https://github.com/ternava/x264/tree/x264-rmv>

## Remained challenges:

- Automation of our approach, while maintaining system validation
- Generalisation of the approach and system validation
- Exploring the benefits on other (non-)functional properties